All Readings:
Introduction to Generative AI

Kindly note that the 30 minutes indicated on the platform considers the time that it may take you to browse through the reading resources provided. The total time required depends on the readings you decide to explore further.

Assembled readings on generative AI:

- Ask a Techspert: What is generative AI?
  https://blog.google/inside-google/googlers/ask-a-techspert/what-is-generative-ai/

- What is generative AI?
  https://www.mckinsey.com/featured-insights/mckinsey-explainers/what-is-generative-ai

- Google Research, 2022 & beyond: Generative models:
  https://ai.googleblog.com/2023/01/google-research-2022-beyond-language.html#GenerativeModels

- Building the most open and innovative AI ecosystem:
  https://cloud.google.com/blog/products/ai-machine-learning/building-an-open-generative-ai-partner-ecosystem

- Generative AI is here. Who Should Control It?
  https://www.nytimes.com/2022/10/21/podcasts/hard-fork-generative-artificial-intelligence.html

- Stanford U & Google's Generative Agents Produce Believable Proxies of Human Behaviors:
  https://syncedreview.com/2023/04/12/stanford-u-googles-generative-agents-produce-believable-proxies-of-human-behaviours/

- Generative AI: Perspectives from Stanford HAI:
  https://hai.stanford.edu/sites/default/files/2023-03/Generative_AI_HAI_Perspectives

- Generative AI at Work:
  https://www.nber.org/system/files/working_papers/w31161/w31161.pdf

- The future of generative AI is niche, not generalized:
  https://www.technologyreview.com/2023/04/27/1072102/the-future-of-generative-ai-is-niche-not-generalized/

- The implications of Generative AI for businesses: https://www2.deloitte.com/us/en/pages/consulting/articles/generative-artificial-intelligence.html

- Proactive Risk Management in Generative AI: https://www2.deloitte.com/us/en/pages/consulting/articles/responsible-use-of-generative-ai.html

- How Generative AI Is Changing Creative Work: https://hbr.org/2022/11/how-generative-ai-is-changing-creative-work

Assembled readings on large language models:

- NLP's ImageNet moment has arrived: https://thegradient.pub/nlp-imagenet/

- LaMDA: our breakthrough conversation technology: https://blog.google/technology/ai/lamda/

- Language Models are Few-Shot Learners: https://proceedings.neurips.cc/paper/2020/file/1457c0d6bfcb4967418bfb8ac142f64a-Paper.pdf

- Introducing Gemini: our largest and most capable AI model: https://blog.google/technology/ai/google-gemini-ai/#sundar-note

- The Power of Scale for Parameter-Efficient Prompt Tuning: https://arxiv.org/pdf/2104.08691.pdf

- Google Research, 2022 & beyond: Language models: https://ai.googleblog.com/2023/01/google-research-2022-beyond-language.html#LanguageModels

- Solving a machine-learning mystery: https://news.mit.edu/2023/large-language-models-in-context-learning-0207

Additional Resources:

- Attention is All You Need: https://research.google/pubs/pub46201/

- Transformer: A Novel Neural Network Architecture for Language Understanding: https://ai.googleblog.com/2017/08/transformer-novel-neural-network.html

- Transformer on Wikipedia:

https://en.wikipedia.org/wiki/Transformer_(machine_learning_model)#:~:text=Transformers%20were%20introduced%20in%202017,allowing%20training%20on%20larger%20datasets.

- What is Temperature in NLP?
  https://lukesalamone.github.io/posts/what-is-temperature/

- Model Garden: https://cloud.google.com/model-garden

- Auto-generated Summaries in Google Docs:
  https://ai.googleblog.com/2022/03/auto-generated-summaries-in-google-docs.html